

## ✓ Simple Logistic Regression

```
#Importing the necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
data = pd.read_csv('/content/drive/MyDrive/ITC Projects/Project_4/Churn_Modelling.csv')
data.head()
print(data.columns)
print(data.info())
print(data.duplicated().sum())
data = data.drop(['RowNumber', 'CustomerId', 'Surname'], axis = 1)
print(data.shape)
```

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
      'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname                10000 non-null  object
3   CreditScore            10000 non-null  int64
4   Geography              10000 non-null  object
5   Gender                 10000 non-null  object
6   Age                   10000 non-null  int64
7   Tenure                 10000 non-null  int64
8   Balance                10000 non-null  float64
9   NumOfProducts          10000 non-null  int64
10  HasCrCard              10000 non-null  int64
11  IsActiveMember         10000 non-null  int64
12  EstimatedSalary        10000 non-null  float64
13  Exited                 10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
None
0
(10000, 11)
```

## ✓ Exploratory Data Analysis

## Univariate Analysis

```
import seaborn as sns
sns.histplot(data['Age'], bins = 20, color = '#D982B5', kde = True)
plt.xlabel('Age')
plt.ylabel('Count')
plt.title('Age Distribution')
plt.show()



sns.histplot(data['CreditScore'], bins = 20, color = '#D982B5', kde = True)
plt.xlabel('Credit Score')
plt.ylabel('Count')
plt.title('Credit Score Distribution')
plt.show()

sns.histplot(data['Gender'], bins = 2, color = '#D982B5')
plt.xlabel('Gender')
plt.ylabel('Count')
plt.title('Gender Distribution')
plt.show()


sns.histplot(data['Geography'], bins = 2, color = '#D982B5')
plt.xlabel('Geography')
plt.ylabel('Count')
plt.title('Geography Distribution')
plt.show()



avg_age_exited = data.groupby('Exited')['Age'].mean()
plt.bar(avg_age_exited.index, avg_age_exited.values, color = '#D982B5')
plt.xlabel('Exited (0 = No, 1 = Yes)')
plt.ylabel('Average Age')
plt.title('Average Age of Customers per Exit Status')
plt.xticks([0, 1], ['0', '1'])
plt.show()


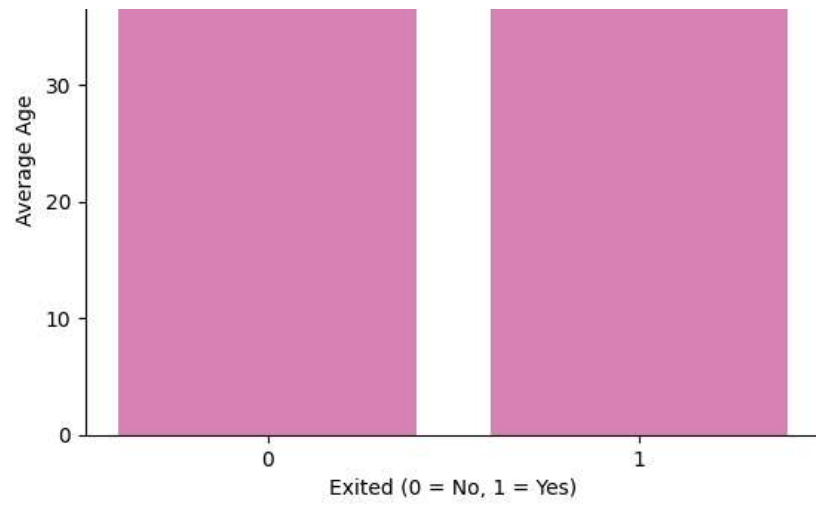
target = data['Exited'].value_counts()
fig1, ax1 = plt.subplots()
ax1.pie(data['Exited'].value_counts(), labels=target.index, autopct='%1.1f%%', colors = ['#D982B5', '#FF9999'], shadow=None)
plt.show()
```



Gender Distribution







```
Count_Exited = data['Exited'].value_counts()
Count_Gender = data['Gender'].value_counts()
Count_Geography = data['Geography'].value_counts()

print(Count_Exited)
print(Count_Gender)
print(Count_Geography)
```

```
↔ Exited
0    7963
1    2037
Name: count, dtype: int64
Gender
Male      5457
Female    4543
Name: count, dtype: int64
Geography
France    5014
Germany   2509
Spain     2477
Name: count, dtype: int64
```

## ▼ Bivariate Analysis

```
sns.barplot(x = 'Exited', y = 'Age', data = data, color = '#967BB6')
plt.xlabel('Exited')
plt.ylabel('Age')
plt.title('Age per Exited')
plt.show()

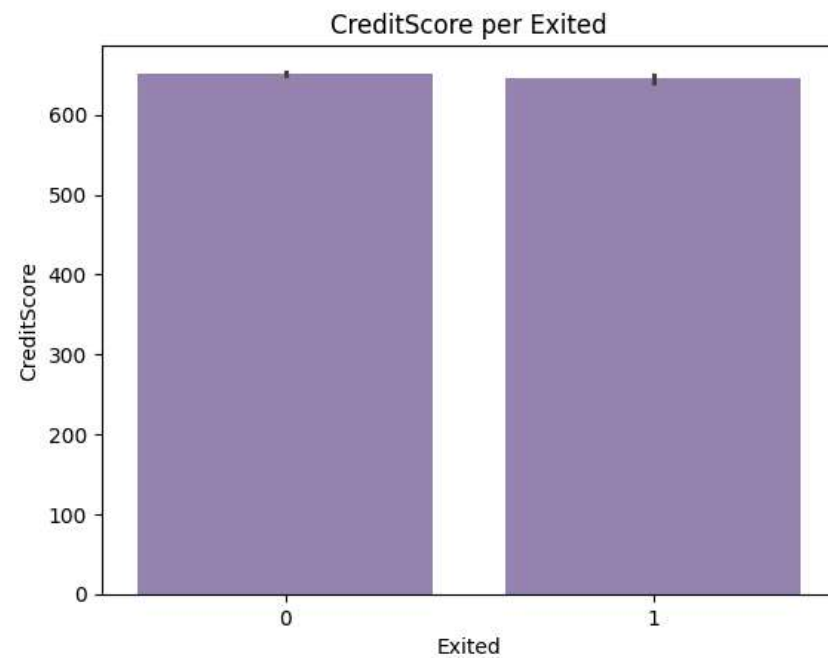
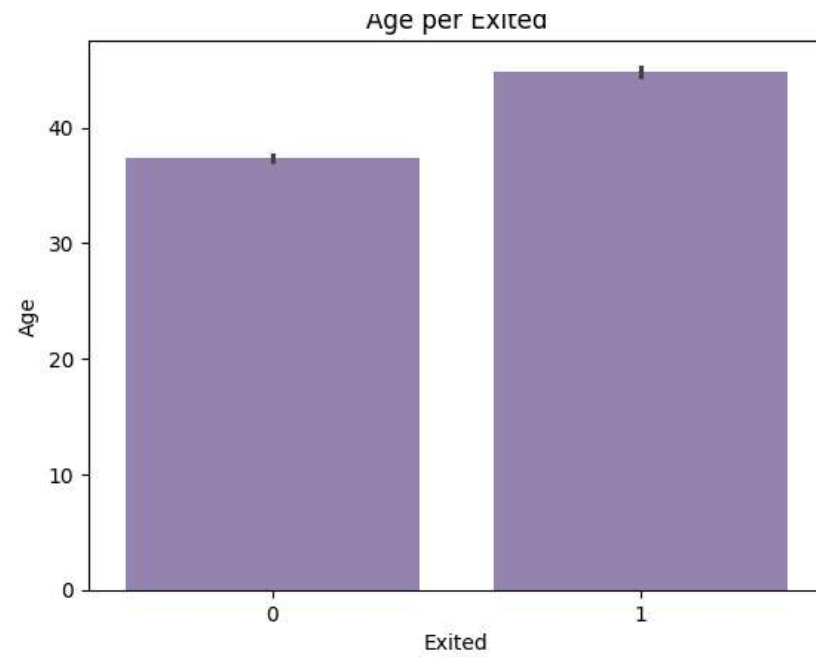
sns.barplot(x = 'Exited', y = 'CreditScore', data = data, color = '#967BB6')
plt.xlabel('Exited')
plt.ylabel('CreditScore')
plt.title('CreditScore per Exited')
plt.show()

sns.barplot(x = 'Exited', y = 'Balance', data = data, color = '#967BB6')
plt.xlabel('Exited')
plt.ylabel('Balance')
plt.title('Balance per Exited')
plt.show()

sns.barplot(x = 'Exited', y = 'NumOfProducts', data = data, color = '#967BB6')
plt.xlabel('Exited')
plt.ylabel('Number of Products')
plt.title('Number of Products per Exited')
plt.show()

sns.barplot(x = 'Exited', y = 'EstimatedSalary', data = data, color = '#967BB6')
```

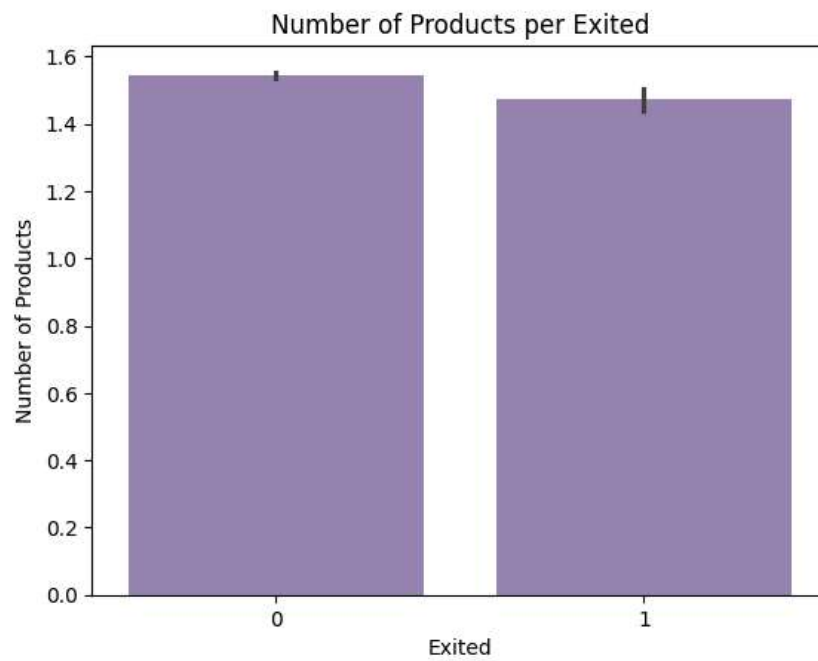
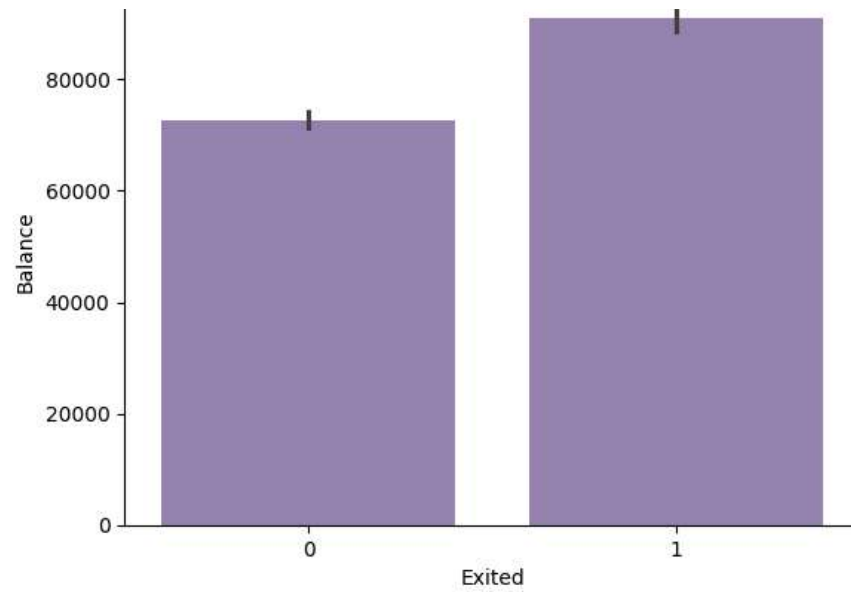
```
plt.xlabel('Exited')  
plt.ylabel('Estimated Salary')  
plt.title('Estimated Salary per Exited')  
plt.show()
```

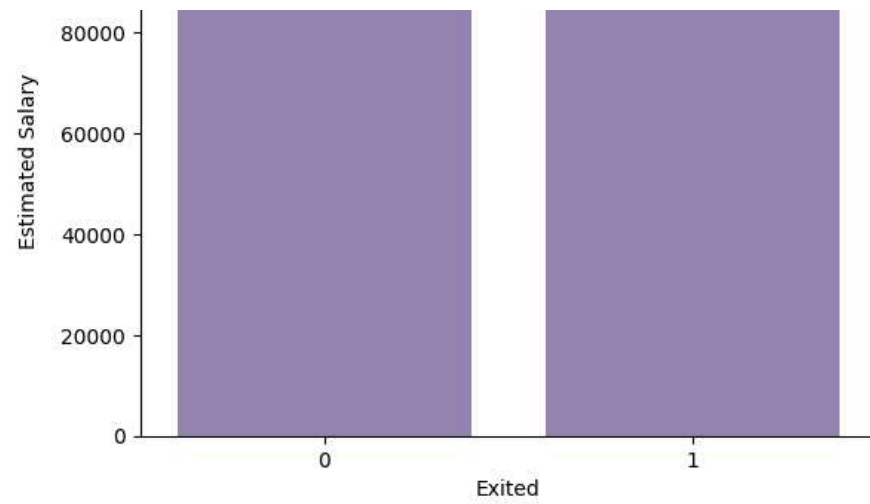


Balance per Exited









```

Avg_creditScore_by_Exited = data.groupby('Exited')['CreditScore'].mean()
Avg_Age_by_Exited = data.groupby('Exited')['Age'].mean()
Avg_Balance_by_Exited = data.groupby('Exited')['Balance'].mean()
Avg_NumOfProducts_by_Exited = data.groupby('Exited')['NumOfProducts'].mean()
Avg_EstimatedSalary_by_Exited = data.groupby('Exited')['EstimatedSalary'].mean()

```

```

print(Avg_creditScore_by_Exited)
print(Avg_Age_by_Exited)
print(Avg_Balance_by_Exited)
print(Avg_NumOfProducts_by_Exited)
print(Avg_EstimatedSalary_by_Exited)

```

```

↗ Exited
0    651.853196
1    645.351497
Name: CreditScore, dtype: float64
Exited
0    37.408389
1    44.837997
Name: Age, dtype: float64
Exited
0    72745.296779
1    91108.539337
Name: Balance, dtype: float64
Exited
0    1.544267
1    1.475209
Name: NumOfProducts, dtype: float64
Exited
0    99738.391772
1    101465.677531
Name: EstimatedSalary, dtype: float64


```

```

data1 = data[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited']]
plt.figure(figsize = (10,6))
sns.heatmap(round(data1.corr(),1), cmap="PuRd", annot =True, linewidths= 0.5)

```

&lt;Axes: &gt;



```

Geography_Encoded= pd.get_dummies(data['Geography'], dtype = int)
Gender_Encoded= pd.get_dummies(data['Gender'], dtype = int)

df = pd.concat([data, Geography_Encoded], axis = 1)
df = pd.concat([df, Gender_Encoded], axis = 1)

df = df.drop(['Geography', 'Gender'], axis = 1)
df.head()

```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	France	Germany	Spain	Female	Male
0	619	42	2	0.00	1	1	1	101348.88	1	1	0	0	1	0
1	608	41	1	83807.86	1	0	1	112542.58	0	0	0	1	1	0
2	502	42	8	159660.80	3	1	0	113931.57	1	1	0	0	1	0
3	699	39	1	0.00	2	0	0	93826.63	0	1	0	0	1	0
4	850	43	2	125510.82	1	1	1	79084.10	0	0	0	1	1	0

Next steps: [View recommended plots](#) [New interactive sheet](#)



```
#MinMaxScaling
data_scaled = (df - df.min())/(df.max()-df.min())
data_scaled
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	France	Germany	Spain	Female	Male
0	0.538	0.324324	0.2	0.000000	0.000000	1.0	1.0	0.506735	1.0	1.0	0.0	0.0	1.0	0.0
1	0.516	0.310811	0.1	0.334031	0.000000	0.0	1.0	0.562709	0.0	0.0	0.0	1.0	1.0	0.0
2	0.304	0.324324	0.8	0.636357	0.666667	1.0	0.0	0.569654	1.0	1.0	0.0	0.0	1.0	0.0
3	0.698	0.283784	0.1	0.000000	0.333333	0.0	0.0	0.469120	0.0	1.0	0.0	0.0	1.0	0.0
4	1.000	0.337838	0.2	0.500246	0.000000	1.0	1.0	0.395400	0.0	0.0	0.0	1.0	1.0	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	0.842	0.283784	0.5	0.000000	0.333333	1.0	0.0	0.481341	0.0	1.0	0.0	0.0	0.0	1.0
9996	0.332	0.229730	1.0	0.228657	0.000000	1.0	1.0	0.508490	0.0	1.0	0.0	0.0	0.0	1.0
9997	0.718	0.243243	0.7	0.000000	0.000000	0.0	1.0	0.210390	1.0	1.0	0.0	0.0	1.0	0.0
9998	0.844	0.324324	0.3	0.299226	0.333333	1.0	0.0	0.464429	1.0	0.0	1.0	0.0	0.0	1.0
9999	0.884	0.135135	0.4	0.518708	0.000000	1.0	0.0	0.190914	0.0	1.0	0.0	0.0	1.0	0.0

10000 rows x 14 columns

Next steps: [View recommended plots](#) [New interactive sheet](#)

```
plt.figure(figsize = (10,6))
sns.heatmap(round(data_scaled.corr(),1), cmap = "PuRd", annot =True, linewidths= 0.5)
```

 <Axes: >


```
data_scaled = data_scaled.drop(['France', 'Male'], axis = 1)
```

```
data_scaled.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   CreditScore     10000 non-null  float64
1   Age             10000 non-null  float64
2   Tenure          10000 non-null  float64
3   Balance         10000 non-null  float64
4   NumOfProducts   10000 non-null  float64
5   HasCrCard       10000 non-null  float64
6   IsActiveMember  10000 non-null  float64
```

```
7 EstimatedSalary 10000 non-null float64
8 Exited          10000 non-null float64
9 Germany         10000 non-null float64
10 Spain          10000 non-null float64
11 Female         10000 non-null float64
dtypes: float64(12)
memory usage: 937.6 KB
```

## ✓ Simple Logistic Regression

## ✓ Useful Functions

```
def sigmoid(h):
    return 1/(1+np.exp(-h))

def cross_entropy(y,p_hat):
    return -(1/len(y))*np.sum(y*np.log(p_hat)+(1-y)*np.log(1-p_hat))

def accuracy (y,y_hat):
    return np.mean(y==y_hat)
```

## ✓ Logistic Regression Class

```
class LogisticRegression():
    def __init__(self, thresh = 0.5):
        # thresh is hyperparameter
        self.thresh = thresh
        self.W = None
        self.b = None

    def fit(self, X, y, eta = 1e-3, epochs = 1e3, show_curve = False):
        epochs = int(epochs)
        N, D = X.shape

        #Initialize Weights and Biases
        self.W = np.random.randn(D)
        self.b = np.random.randn(1)

        J = np.zeros(epochs)

        #SGD
        for epoch in range(epochs):
            p_hat = self.__forward__(X)
            J[epoch] = cross_entropy(y, p_hat)
```

```

#Weight Update Rules
self.W -= eta*(1/N)*X.T@(p_hat-y)
self.b -= eta*(1/N)*np.sum(p_hat-y)

if show_curve:
    plt.figure()
    plt.plot(J)
    plt.xlabel('epochs')
    plt.ylabel('$\mathcal{J}$')
    plt.title("Training Curve")

def __forward__(self, X):
    return sigmoid(X@self.W +self.b)

def predict(self, X):
    return (self.__forward__(X) >= self.thresh).astype(np.int32)

```

## ✓ Train Test Split

```

#With Scaling

train_ratio = 0.8

train_data_log_reg = data_scaled.sample(frac=train_ratio, random_state=42)
test_data_log_reg = data_scaled.drop(train_data_log_reg.index)

print(train_data_log_reg.shape)
print(test_data_log_reg.shape)

y_train = train_data_log_reg['Exited']
X_train = train_data_log_reg.drop(['Exited'], axis = 1)
y_test = test_data_log_reg['Exited']
X_test = test_data_log_reg.drop(['Exited'], axis = 1)

print(y_train.shape)
print(X_train.shape)
print(y_test.shape)
print(X_test.shape)

# converting X_train, X_test, y_train, y_test to NumPy arrays
X_train = np.array(X_train)
X_test = np.array(X_test)
y_train = np.array(y_train)

```



```
y_test = np.array(y_test)
y_train = y_train.astype(int)
```

```
(8000, 12)
(2000, 12)
(8000,)
(8000, 11)
(2000,)
(2000, 11)
```


## ✓ Implementation of Binary Logistic Regression

```
LR = LogisticRegression(thresh = 0.51)
LR.fit(X_train, y_train, eta=9e-2, epochs=5e3, show_curve=True)

y_train_hat = LR.predict(X_train)

print(f"Training Accuracy: {accuracy(y_train, y_train_hat): 0.4f}")
print(LR.W, LR.b)
```

```
Training Accuracy: 0.8114
[-0.58028663  3.37011577 -0.17381636  0.51848868 -0.43894914 -0.13440265
 -0.92612008 -0.02721149  0.73951948 -0.00691515  0.48093174] [-1.95364916]
```




```
y_test_hat = LR.predict(X_test)
print(f"Test Accuracy: {accuracy(y_test, y_test_hat): 0.4f}")
```

➦ Test Accuracy: 0.8190

```
plt.figure(figsize=(10,7))
y_actu = pd.Series(y_test, name='Actual')
y_pred = pd.Series(y_test_hat, name='Predicted')
cm = pd.crosstab(y_actu, y_pred)
ax = sns.heatmap(cm, annot=True, fmt="d", linewidths=.5)
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

➦ Text(0.5, 47.72222222222222, 'Predicted label')



## ✓ Artificial Neural Network with Variable Architecture

### ✓ Useful Activation Functions

```
# Activations

def linear(H):
    return H

def ReLU(H):
    return H*(H>0)

def softmax(H):
    eH=np.exp(H)
    return eH/eH.sum(axis=1, keepdims=True)

#Misc

def one_hot(y):
    N=len(y)
    K=len(set(y))
    Y = np.zeros((N,K))

    for i in range(N):
        Y[i,y[i]]=1

    return Y
```

### ✓ Useful Loss Functions

```
#Loss Functions

def cross_entropy(Y, P_hat):
    return -(1/len(Y))*np.sum(Y*np.log(P_hat))

def OLS(Y, Y_hat):
    return (1/(2*len(Y)))*np.sum((Y-Y_hat)**2)
```

## ✓ Derivatives

```
def derivative(Z, a):

    if a == linear:
        return 1
    elif a == sigmoid:
        return Z*(1-Z)
    elif a==np.tanh:
        return 1-Z*Z
    elif a==ReLU:
        return (Z>0).astype(int)

    else:
        ValueError("Unknown Activation")
```

## ✓ Useful Metrics

```
def accuracy(y,y_hat):
    return np.mean(y==y_hat)

def R2(y,y_hat):
    return 1-np.sum((y-y_hat)**2)/np.sum((y - y.mean())**2)
```

## ✓ ANN Class

```
class ANN():

    def __init__(self, architecture, activations=None, mode=0): #architecture--No. of Neurons in the hidden layers
        self.mode=mode #mode: 0 → Classification (default, uses softmax in the output layer). 1 → Regression (uses linear activation in the output layer).
        self.architecture = architecture
        self.activations = activations
        self.L = len(architecture)+1          #self.L: Total number of layers = hidden layers + output layer.

    def fit (self, X, y, eta=1e-3, epochs=1e3,show_curve=False):

        epochs= int(epochs)
        if self.mode:
            Y=y
            K=1
        else:
            Y = one_hot(y)
```